

Citation Integrity and Claim-Citation Audit

Topic: Token-Level Uncertainty as a Predictor of Hallucination in Large Language Models

This audit checks both reference metadata and the major claims made in the supplied AI-assisted paper.

Reference-level audit

Reference	Status	Action
Fadeeva et al. (2024)	Verified	Authoritative ACL record; included in verified collection.
Kossen et al. (2024)	Verified	Authoritative arXiv record; included in verified collection.
Shelmanov et al. (2025)	Verified / updated	Authoritative ACL record; EMNLP 2025 metadata used.
Villani, M. J. (n.d.)	Needs correction / verification	Incomplete scholarly metadata; excluded from verified count.

Claim-level audit

Area	Status	Finding
CCP / token-level fact checking	Verified	Supported by Fadeeva et al. (2024).
Semantic entropy / SEP	Verified	Supported by Farquhar et al. (2024) and Kossen et al. (2024).
Pre-trained UQ heads	Verified	Supported by Shelmanov et al. (2025).
Villani entropy-fingerprint claims	Needs verification	Source metadata in supplied paper is incomplete.
CES and formal error-guarantee claims	Needs verification	Dependent on the incomplete Villani reference.

Repository verification rule

The repository contains 24 verified scholarly papers, so the minimum 20 is met without relying on the uncertain Villani reference. Third-party research-paper PDFs are not bundled. The repository links to authoritative records instead.